


Balanced Trees


Tree Insertion Sorting

- Construct a binary search tree for all keys, an inorder traversal will visit the elements in sorted order.

| | | | | | | |
|---|---|---|---|---|---|---|
| 6 | 2 | 8 | 5 | 1 | 3 | 4 |
|---|---|---|---|---|---|---|


inorder
traversal 

| | | | | | | |
|---|---|---|---|---|---|---|
| 1 | 2 | 3 | 4 | 5 | 6 | 8 |
|---|---|---|---|---|---|---|




Normal Binary Search Tree

A randomly generated 500-node tree



An Unbalanced BST

- After many *insert* and *delete*, we could end up with an **unbalanced** BST.




Can We Do Better?

- **Balance** is important to ensure that the tree does not degenerate into an unbalanced tree.
- Difficult as balancing a tree often requires the structure of the tree to be changed (take longer on average for updates).


Skewed Binary Search Tree

- Recall that this is also a BST, although it is a *skewed* one.



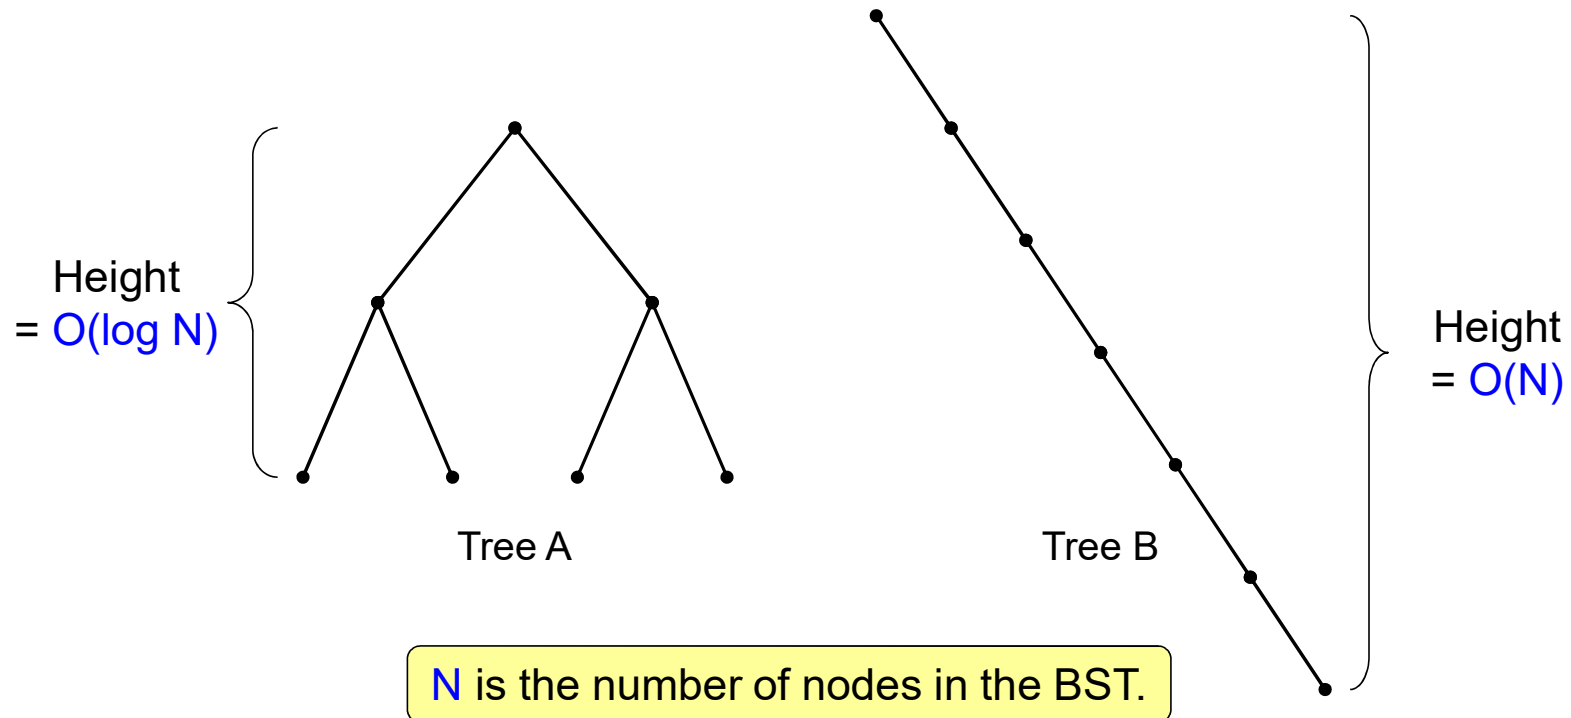
BST Operations

- In fact, the functions **FindNode**, **InsertNode**, and **DeleteNode** all run in $O(H)$ time, where H is the height of the BST.



$O(\log N)$ vs $O(N)$


- Therefore, we like tree A more than tree B, although both trees have 7 nodes.



Balanced Binary Trees

- A binary tree is said to be **balanced** if
 - at each node, the **height** of its **left** and **right** subtrees differ by **at most one**.
- Examples


• The **height** of n_i is the length of the longest path from n_i to a **leaf**.



Unbalanced Binary Trees

- An empty tree has **height -1**


- Examples of **un**balanced binary trees



Tree-balancing Strategy


- We want to avoid the $O(N)$ time complexity associated with unbalanced BSTs.
- The idea is to keep a BST balanced as we build it.
 - During insertion, *rearrange* the nodes in the BST whenever it becomes out of balance.

Balance Factors



- To keep track of whether a BST is balanced, we associate with each node a **balance factor**, which is


$$\text{Height}(\text{right subtree}) - \text{Height}(\text{left subtree})$$




Balance Factors

Height(right subtree) - Height(left subtree)

Node
Height





Balance
Factor



Simple cases


Original blue nodes
in a balance tree.

Addition of new node
that makes the
tree unbalance.






More general cases

Balance factor
before addition.

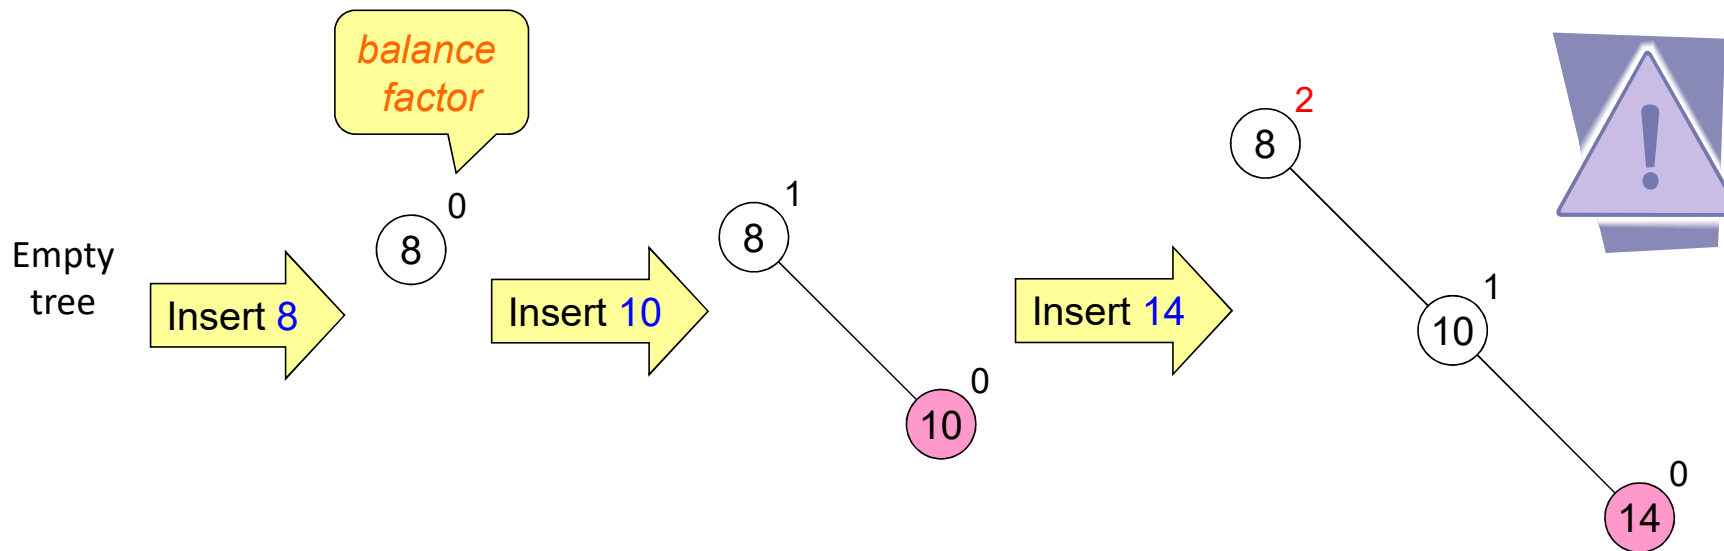


Addition of new node
that makes the
tree unbalance.




Example

- For simplicity, we use numbers as keys.
- From an empty BST, insert the keys 8, 10, and 14.



Fixing the Imbalance


- To fix the imbalance, we *rearrange* the nodes.



1. Node 10 moves *up* to become the root.
2. Node 8 moves *down* to become the child of node 10.

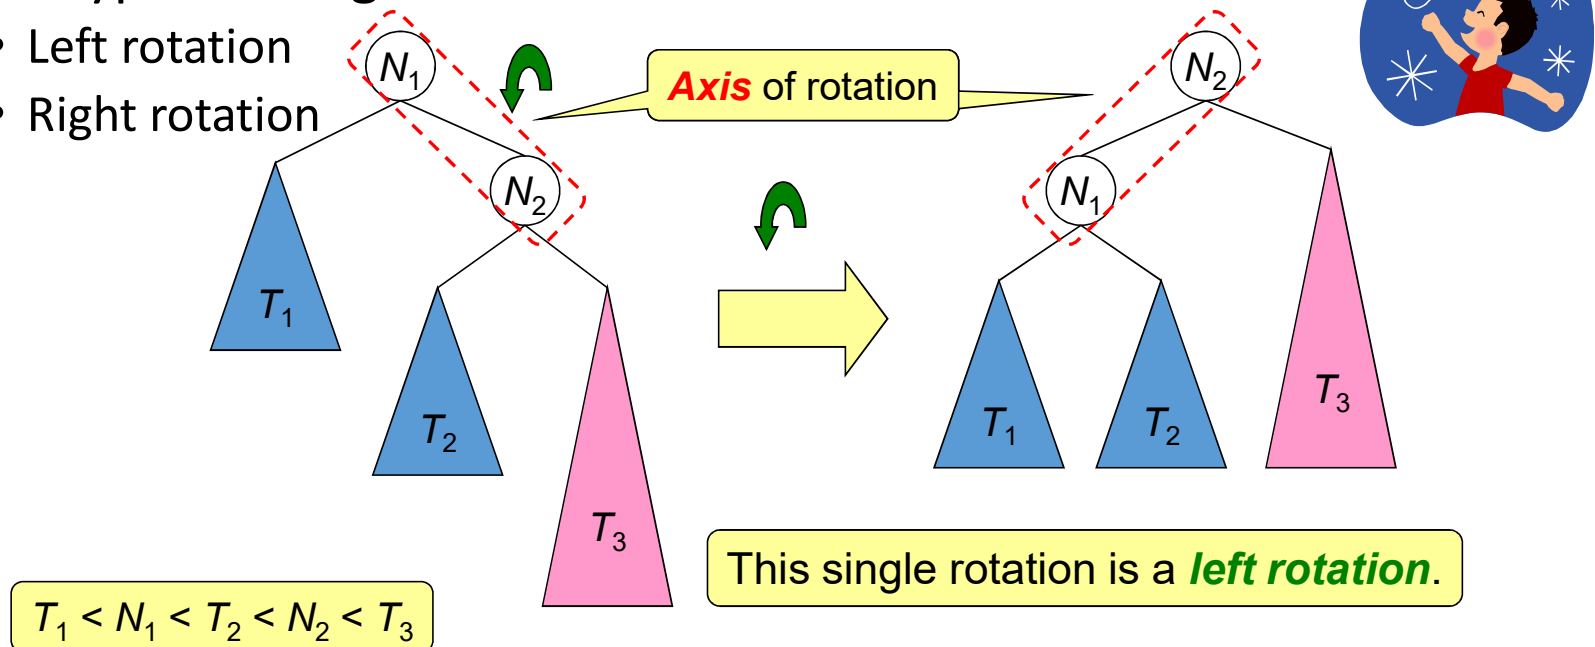
Insertion Causes Imbalance




- In general, such kind of imbalance occurs when a node is inserted into the *right* subtree of the *right* child.



Single Rotation

- To fix such imbalance, we perform an operation called a **single rotation**.
- Two types of single rotation
 - Left rotation
 - Right rotation

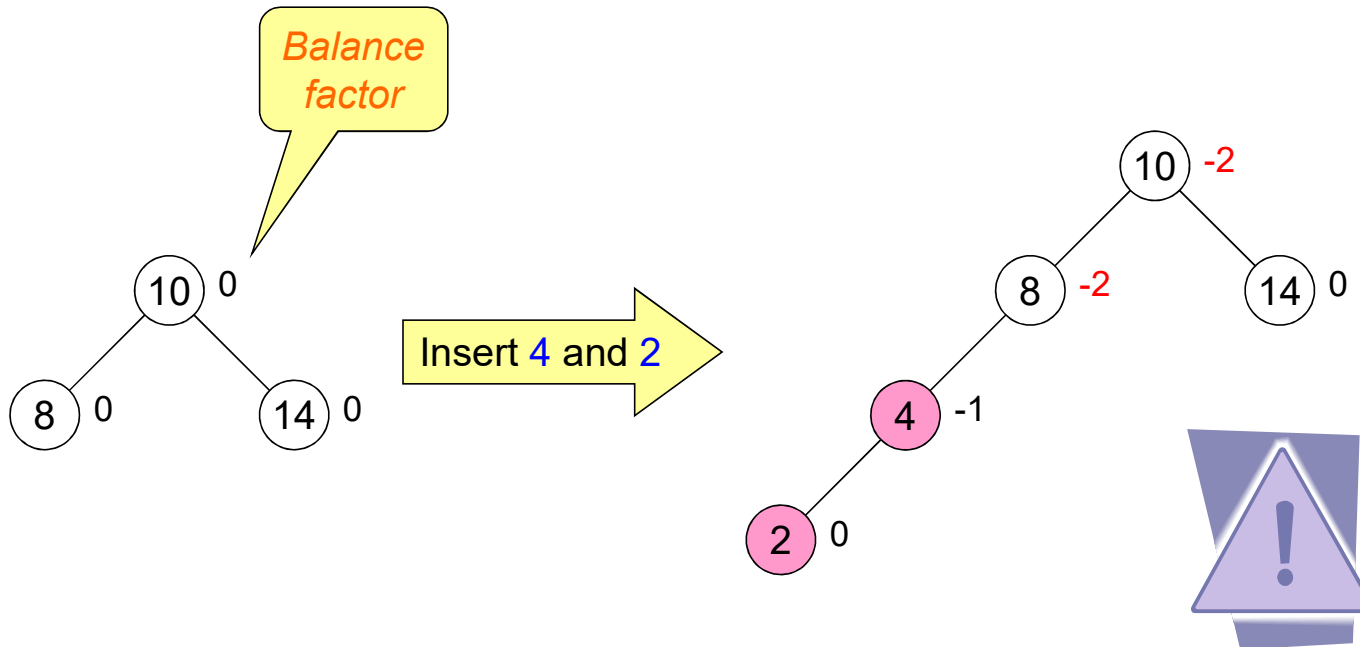




$$T_1 < N_1 < T_2 < N_2 < T_3$$

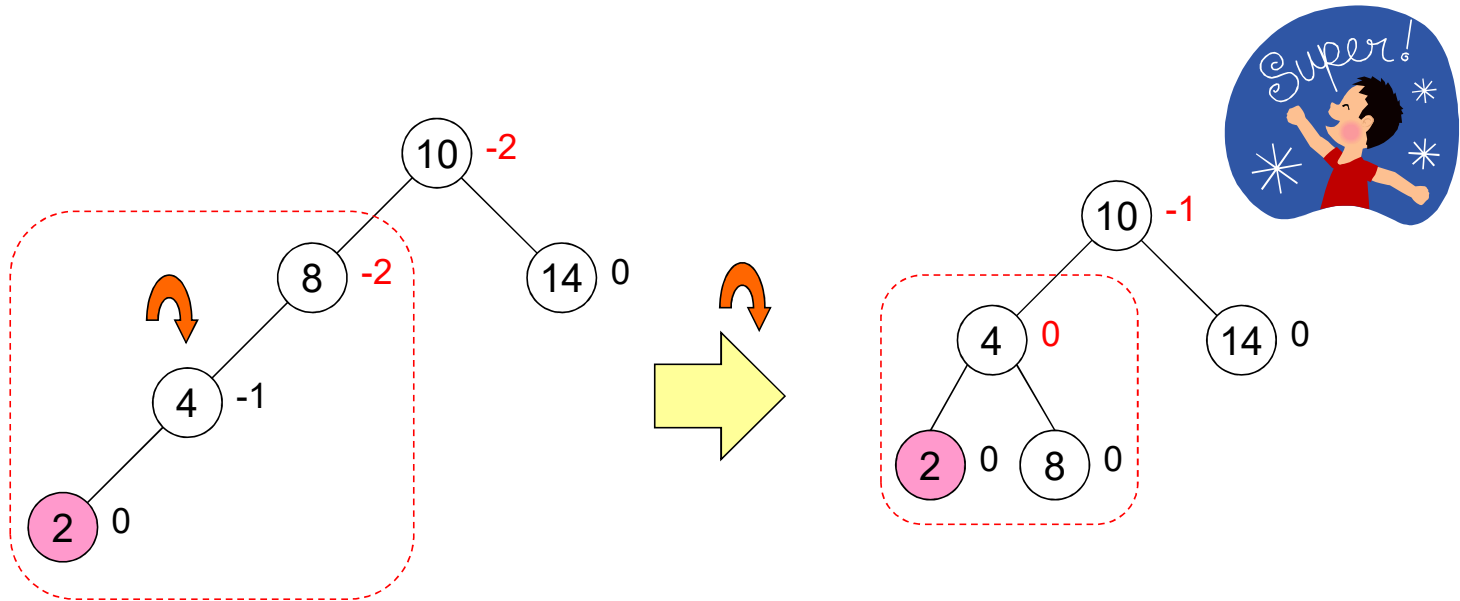
Example

- We continue to insert the keys 4 and 2 into our rotated BST.



Single Rotation

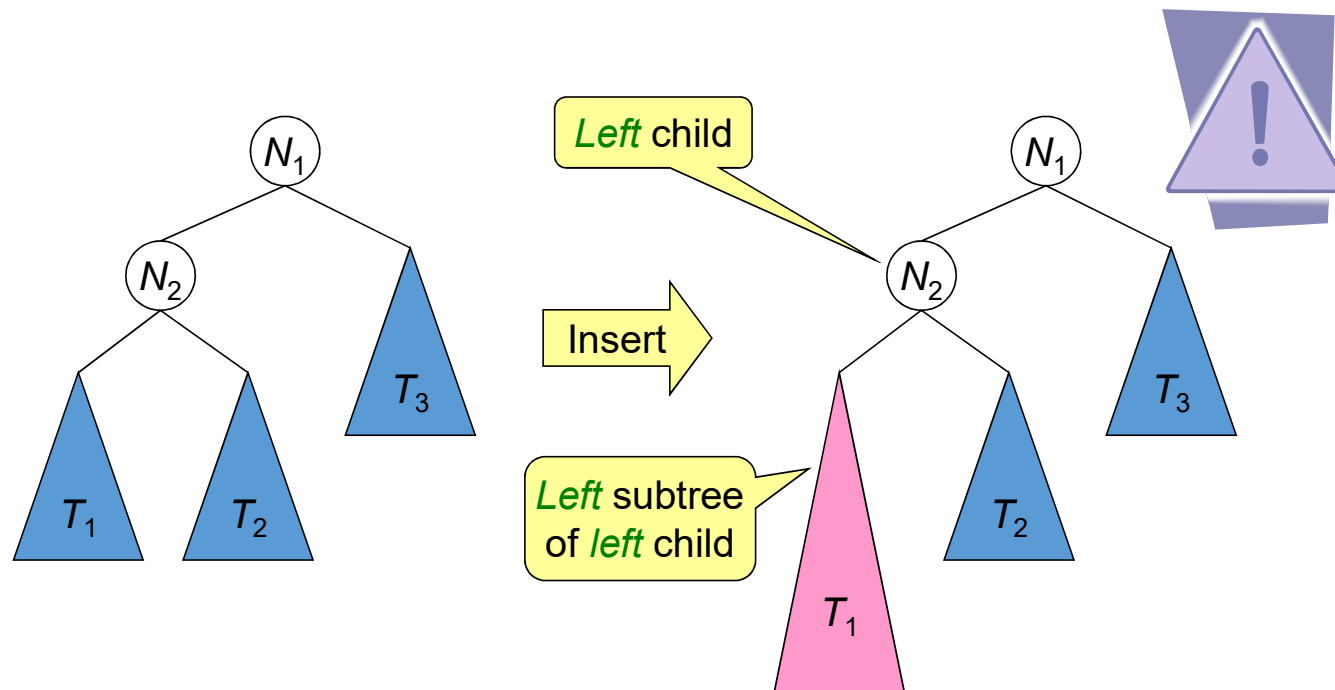
- To fix the imbalance, we perform a *right rotation*.



1. Node 4 moves *up* to become the root of the subtree.
2. Node 8 moves *down* to become the child of node 4.

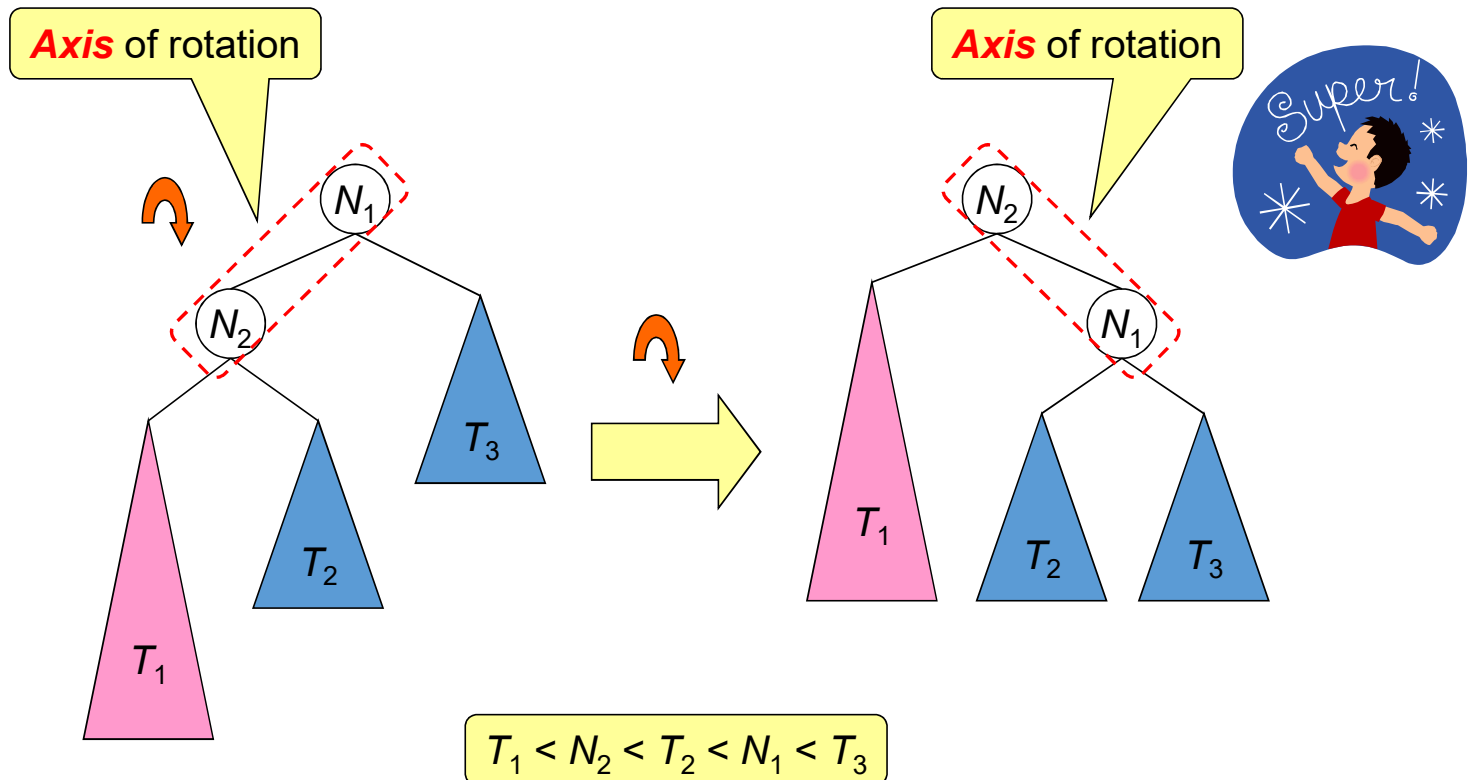
Insertion Causes Imbalance



- In general, such kind of imbalance occurs when a node is inserted into the *left* subtree of the *left* child.






Single Rotation

- A **right rotation** can fix such imbalance.





Axis of rotation




Summary So Far (Single Rotations)

Case 1:

Imbalance occurs when a node is inserted in the *right* subtree of the *right* child.


Solution: *Left* rotation



Case 2:

Imbalance occurs when a node is inserted in the *left* subtree of the *left* child.




Solution: *Right* rotation



Both *left* and *right* rotations are single rotations.

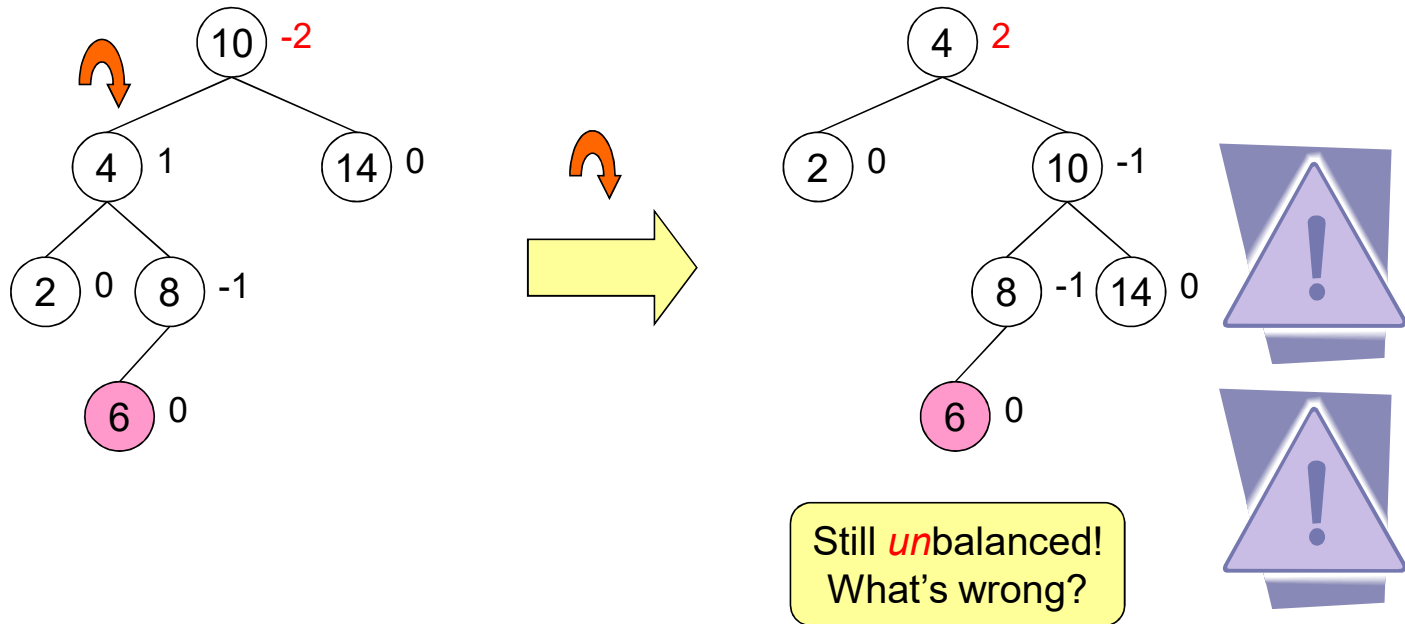
Example

- Next, we insert the key 6.



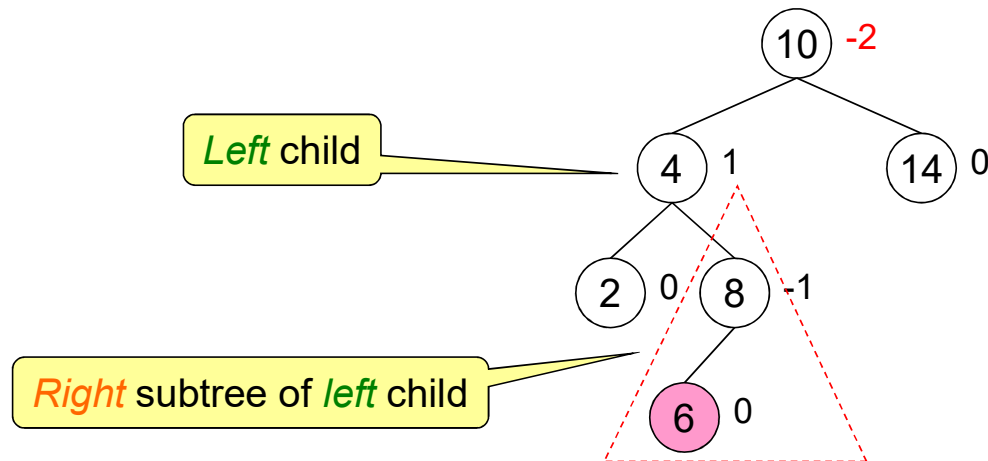
Single Rotation Not Working

- Let's do a *right* rotation along the 10-4 axis.



Single Rotation Not Working

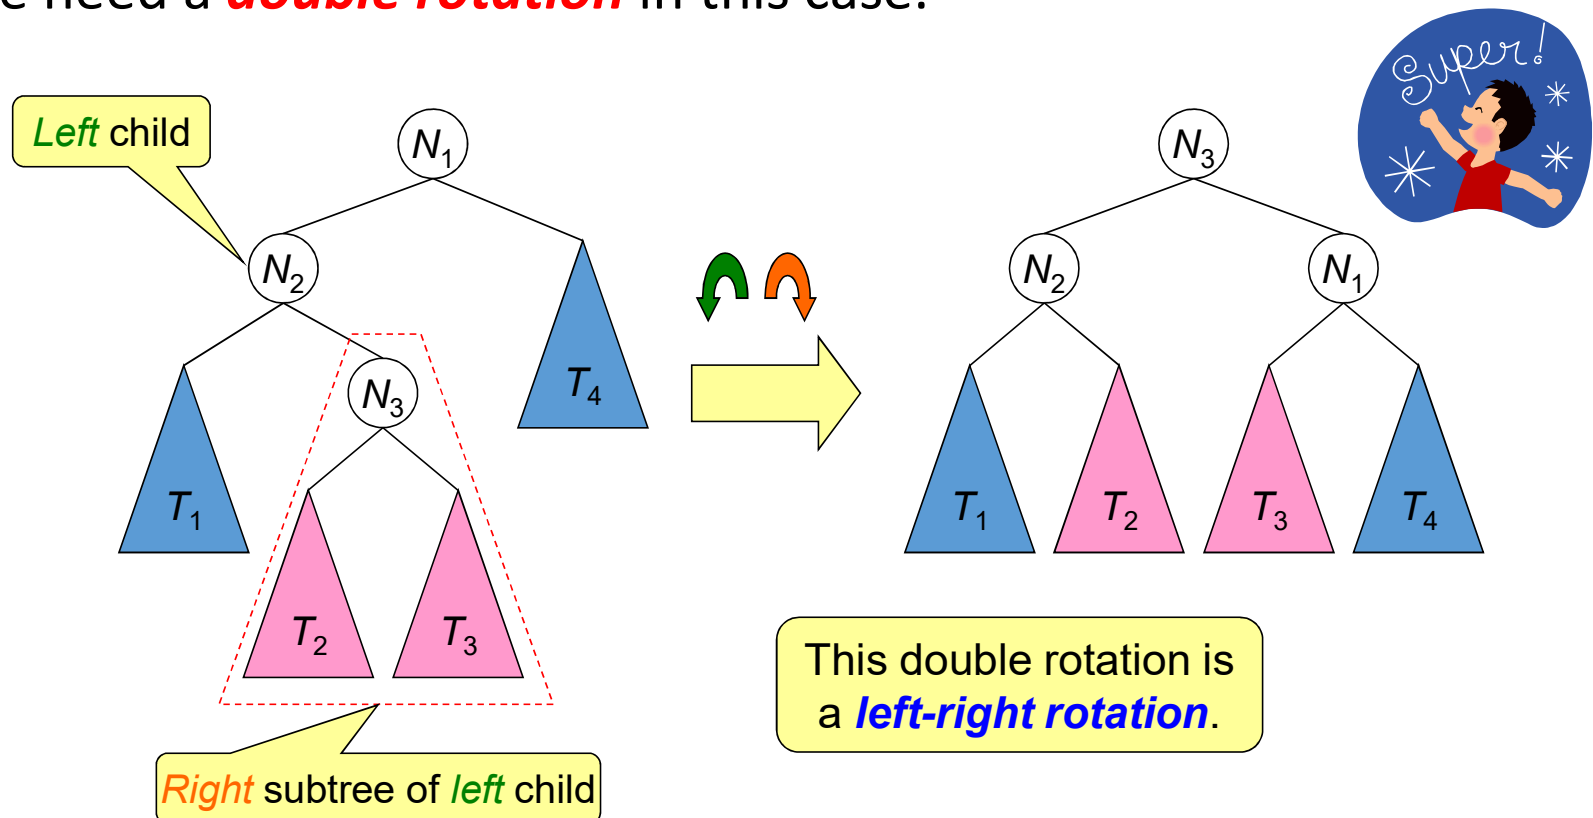
- The problem is, this imbalance occurs when a node is inserted into the *right* subtree of the *left* child.



- Single rotations *cannot* fix such imbalance.


Double Rotation

- We need a **double rotation** in this case.



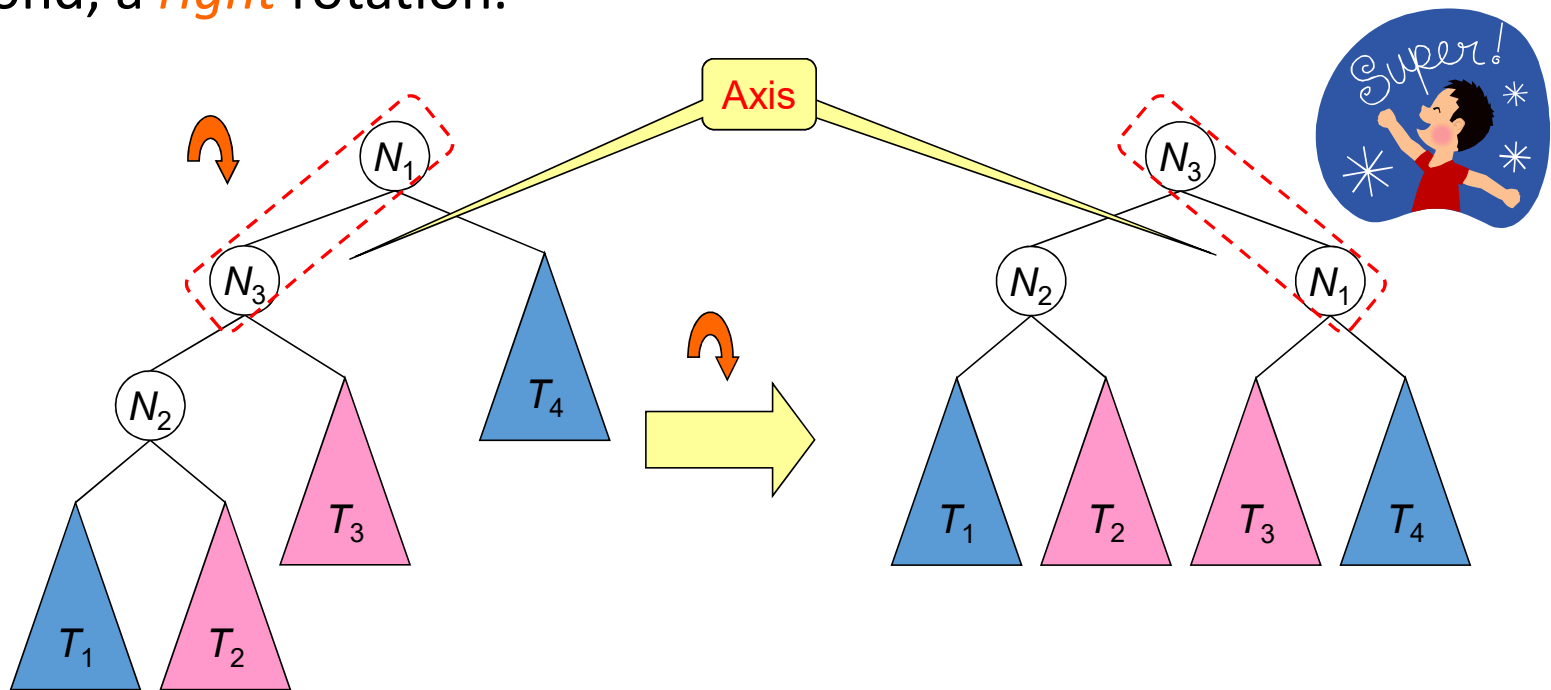
Double Rotation: Step 1

- First, a *left* rotation on the *left* subtree.



Double Rotation: Step 2


- Second, a *right* rotation.






Double Rotation

- A **left-right rotation** is simply a **left** rotation (on the **left** subtree) followed by a **right** rotation.

A **left-right** rotation fixes the imbalance caused by an insertion into the **right** subtree of the **left** child.

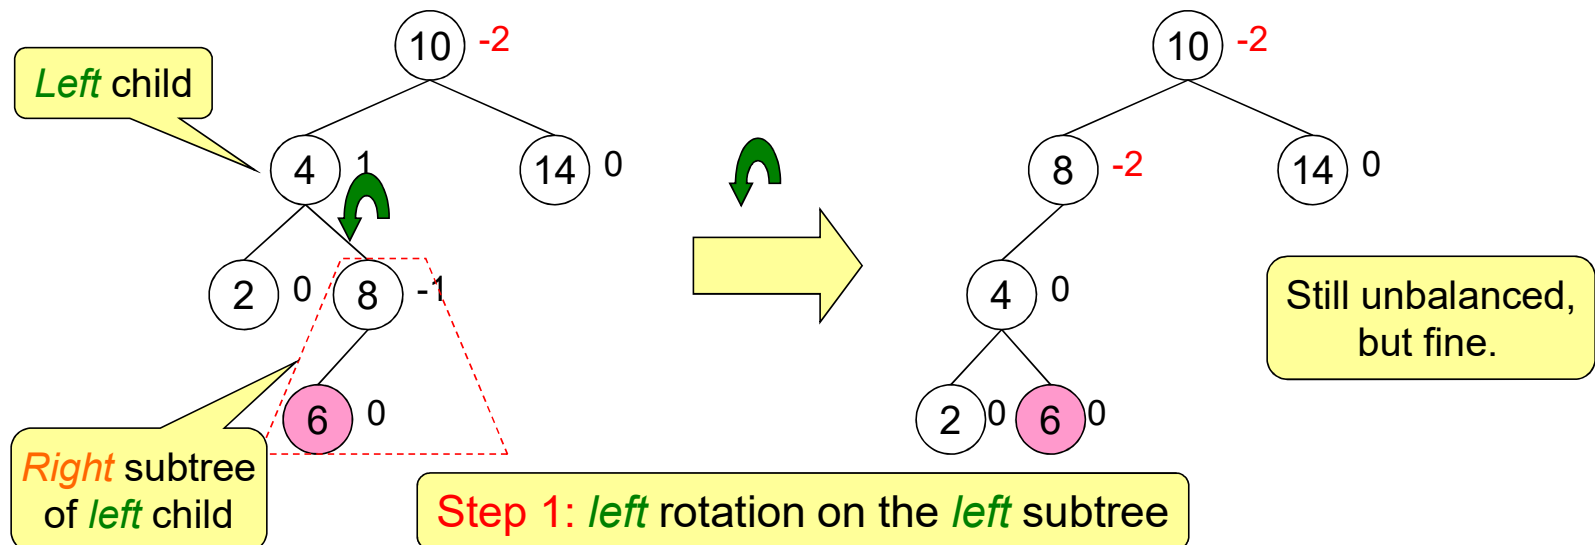




This double rotation is
a **left-right rotation**.


Example

- Let's go back to our example where key **6** is just inserted into the *right* subtree of the *left* child.
- We need to perform a *left-right* rotation.


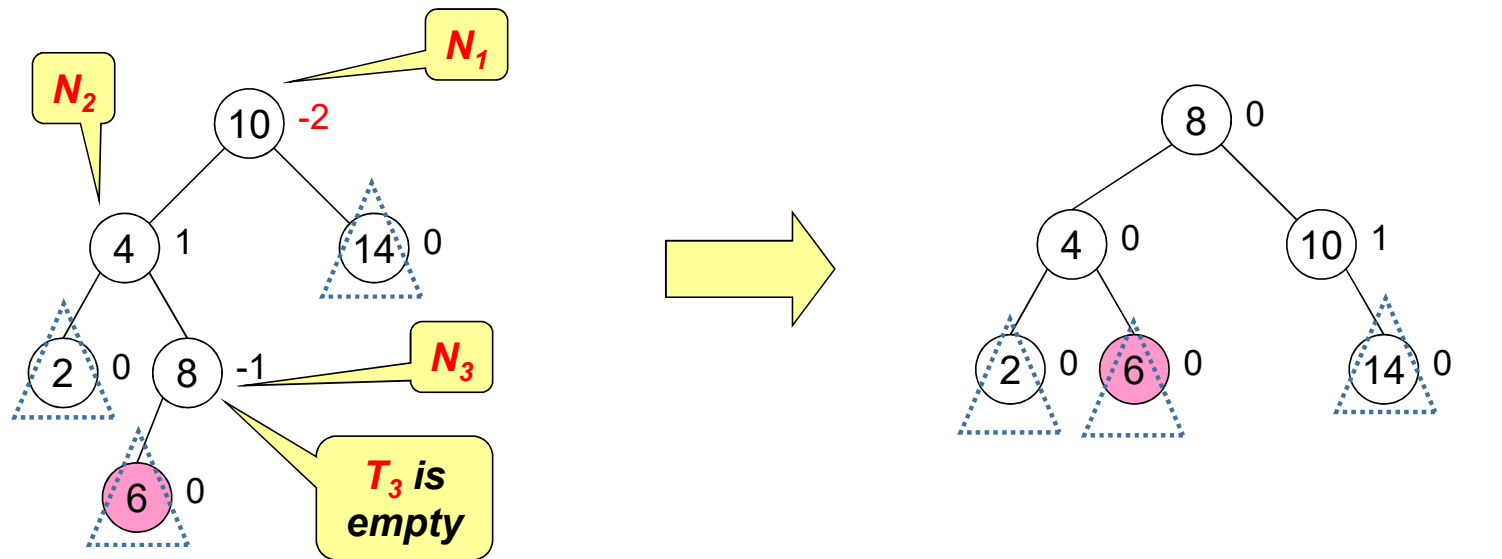


Example

- Continue with the *left-right* rotation.




Step 2: *right* rotation



This double rotation is a *left-right rotation*.

Summary So Far

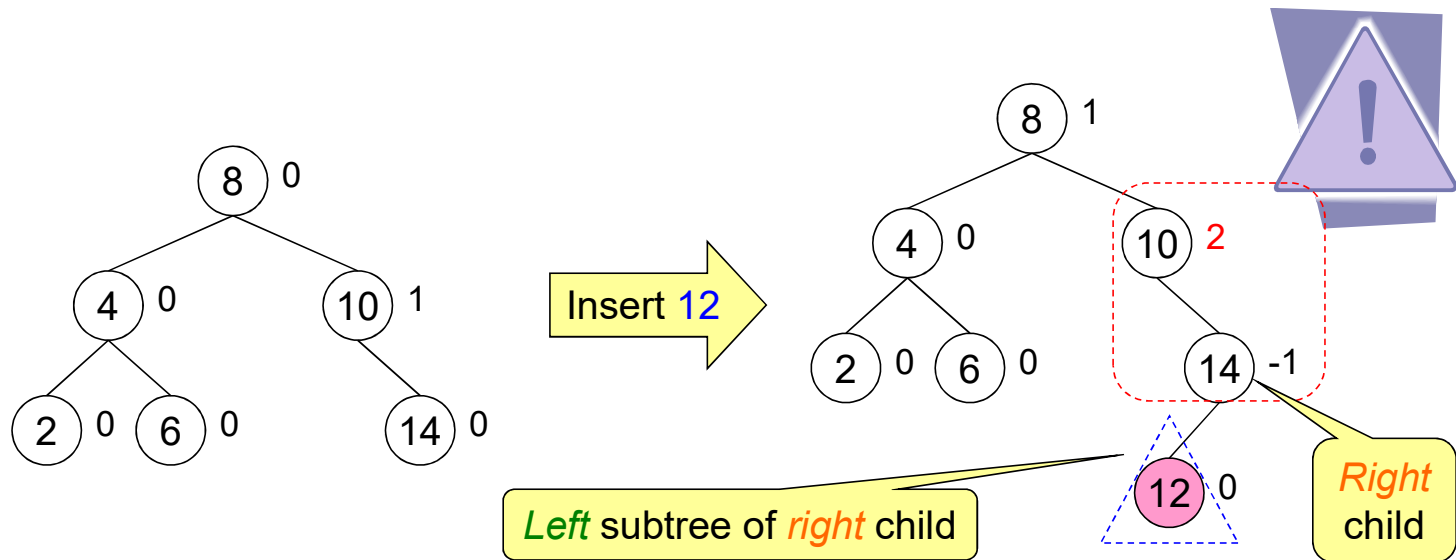
- Single rotations
 - *Left* rotations
 - *Right* rotations
- Double rotations
 - *Left-right* rotations



- As expected, there is a kind of double rotation called *right-left* rotation.

Example

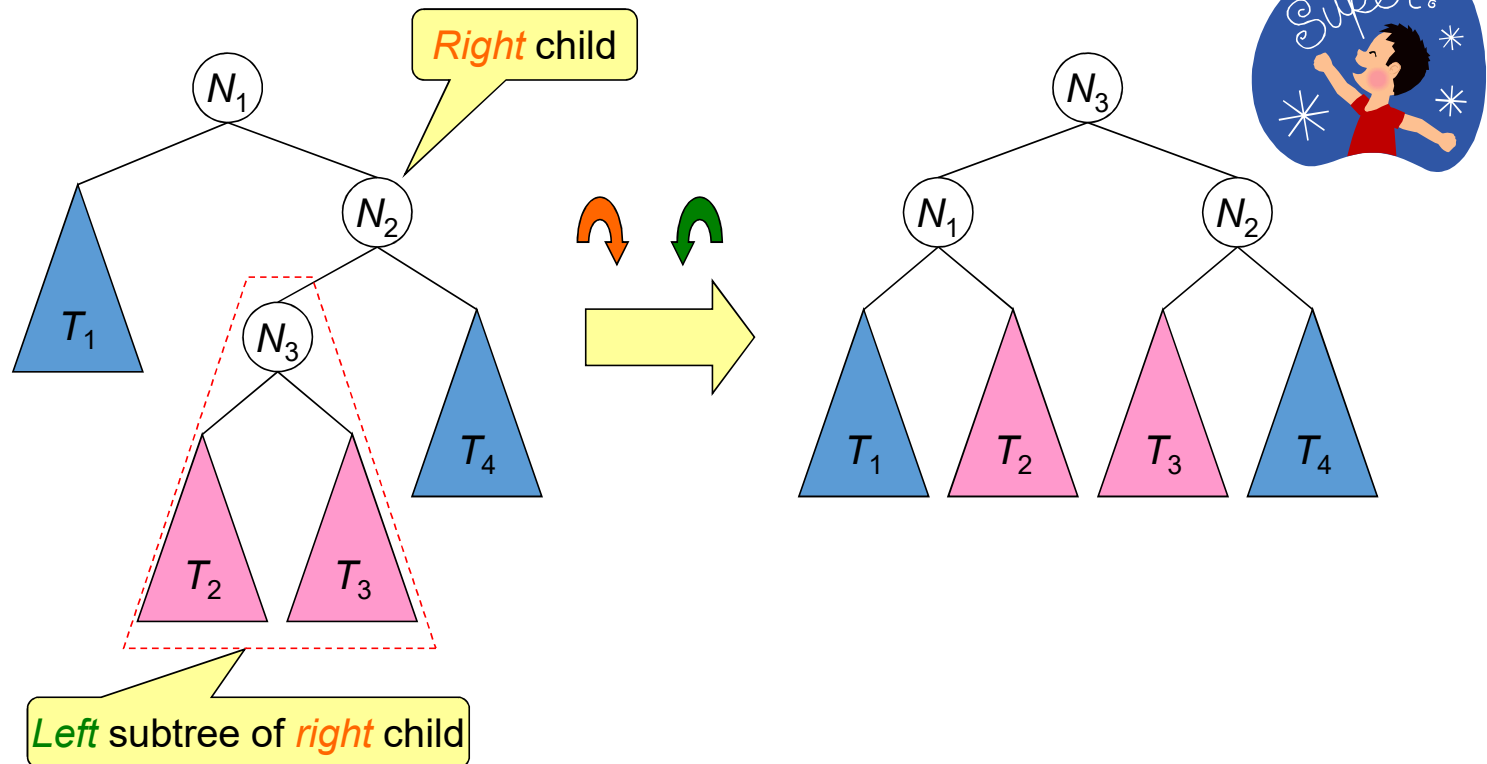
- To illustrate the idea, we insert the key 12.



Such imbalance occurs when a node is inserted into the *left* subtree of the *right* child.

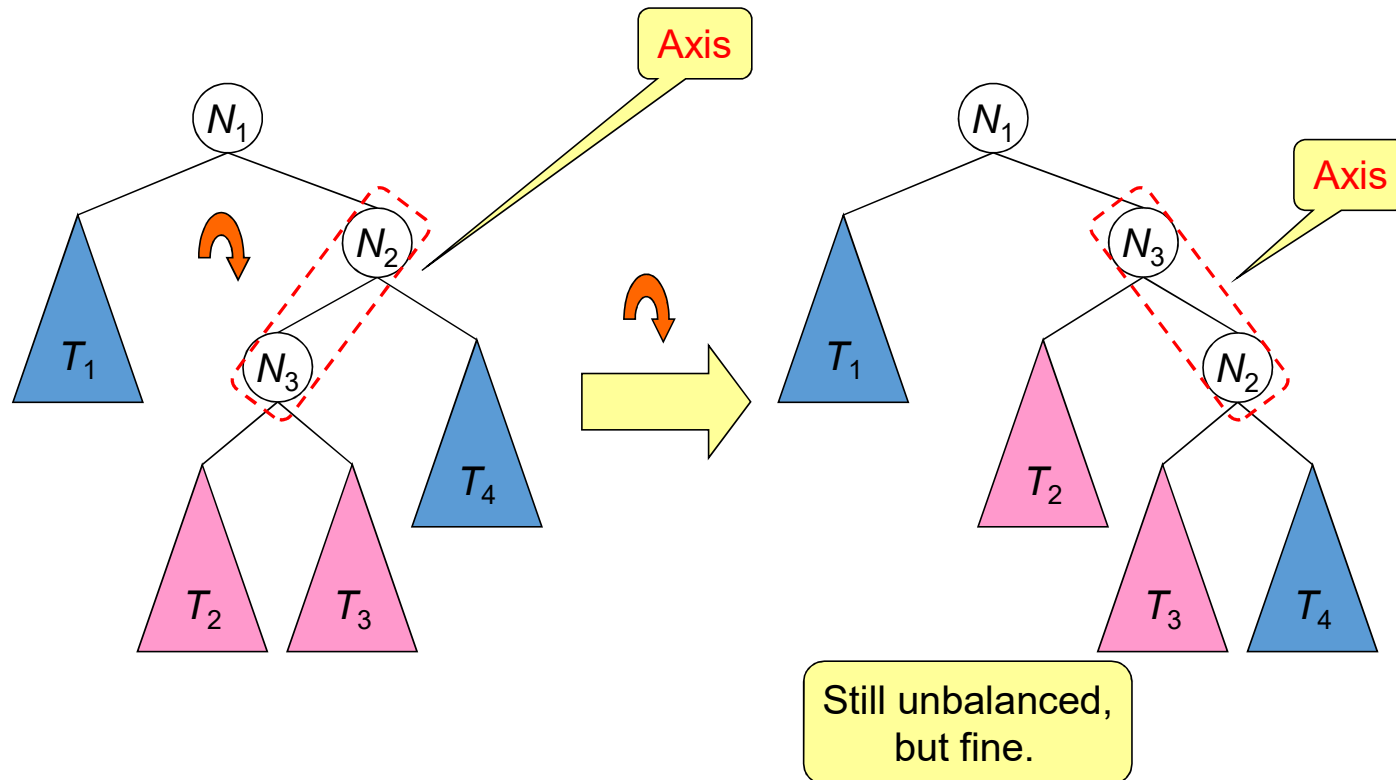
Double Rotation

- We need a **right-left rotation** in this case.



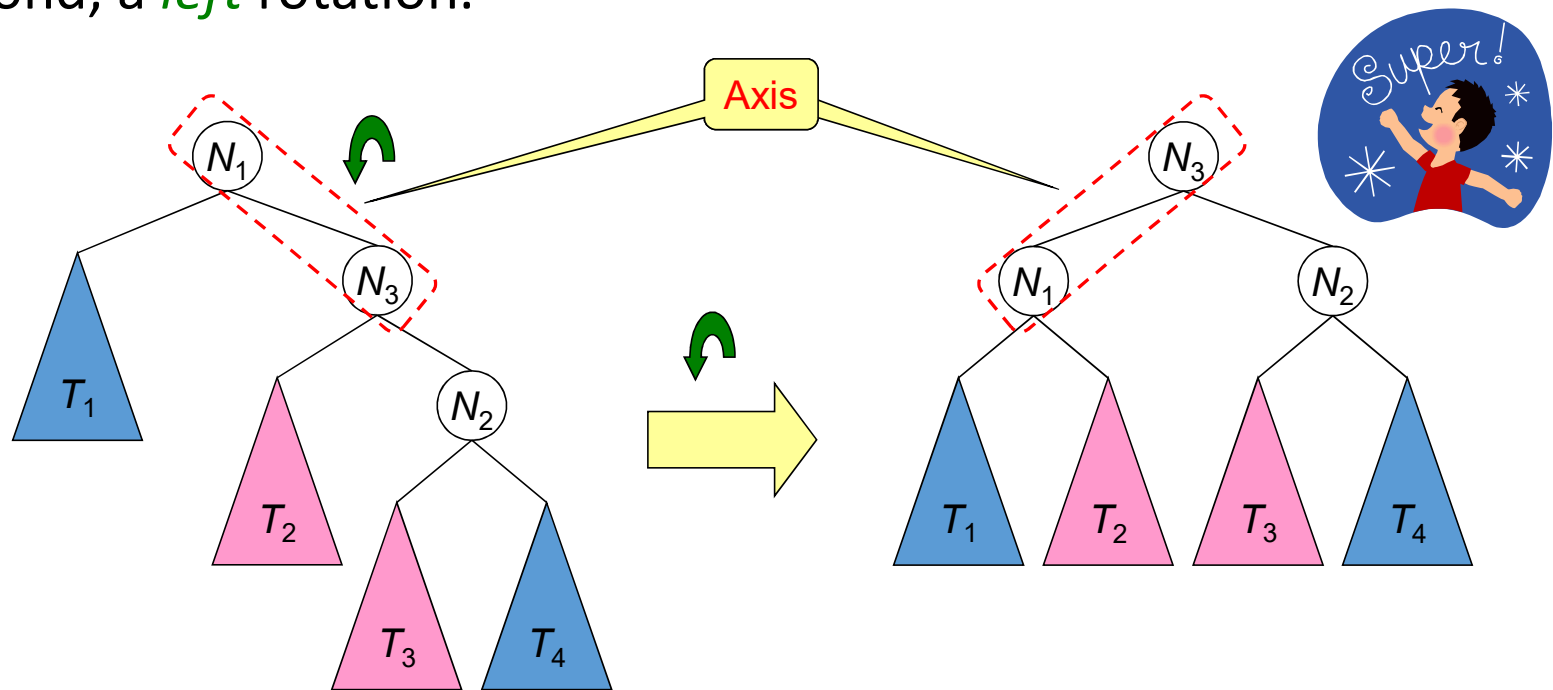
Double Rotation: Step 1

- First, a *right* rotation on the *right* subtree.




Double Rotation: Step 2

- Second, a *left* rotation.



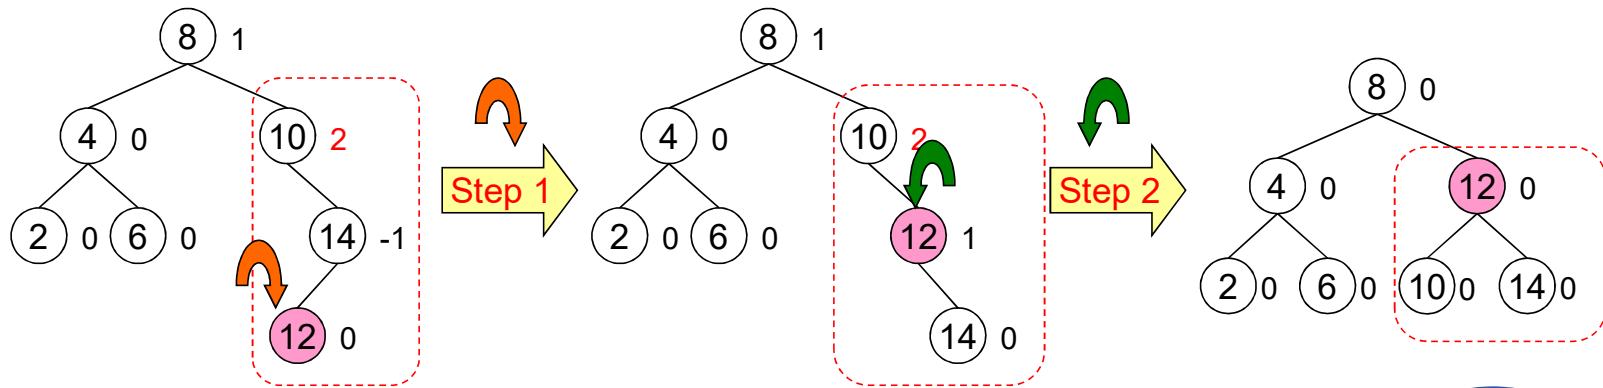
Double Rotation

- A **right-left rotation** is simply a **right** rotation (on the **right** subtree) followed by a **left** rotation.



Example


- Do a *right-left* rotation after inserting key 12.



Summary of Imbalance


Case 1: insertion to *right* subtree of *right* child

Solution: *Left* rotation




Case 2: insertion to *left* subtree of *left* child

Solution: *Right* rotation




Case 3: insertion to *right* subtree of *left* child

Solution: *Left-right* rotation



Case 4: insertion to *left* subtree of *right* child

Solution: *Right-left* rotation



Steps of Fixing Imbalance

- Find the lowest node that is imbalanced, i.e., the lowest node that has a balanced factor ≤ -2 or ≥ 2 , after an insertion
- Determine which case
- Perform the rotation accordingly

AVL Trees

Observation

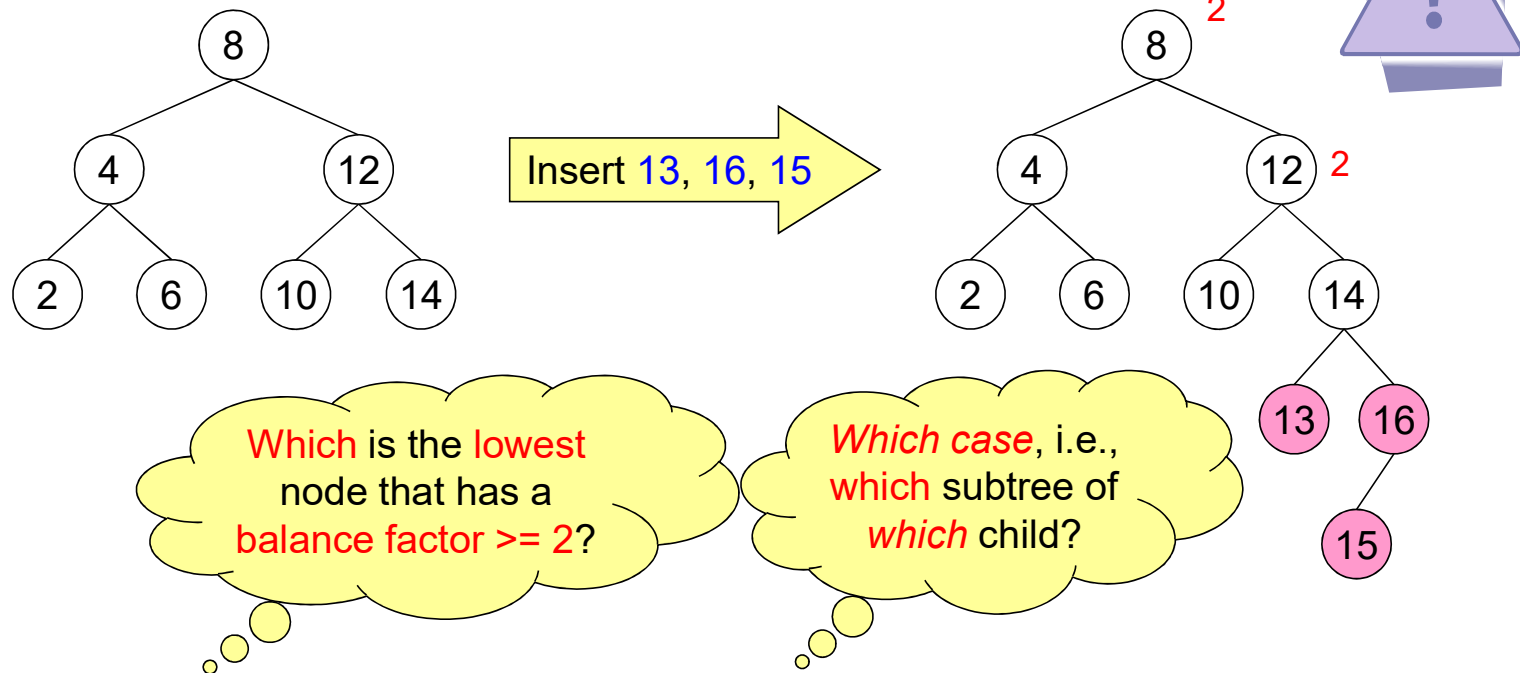
An imbalance caused by insertion can *always* be fixed by performing *one* operation, either a *single* or *double* rotation.

- A balanced tree maintained by single or double rotations is called an *AVL tree*.

AVL stands for **A**del'son-**V**el'skii and **L**andis, the two Russian mathematicians who invented this algorithm.


Example

- To complete the example, we insert the key 13, 16, and 15.

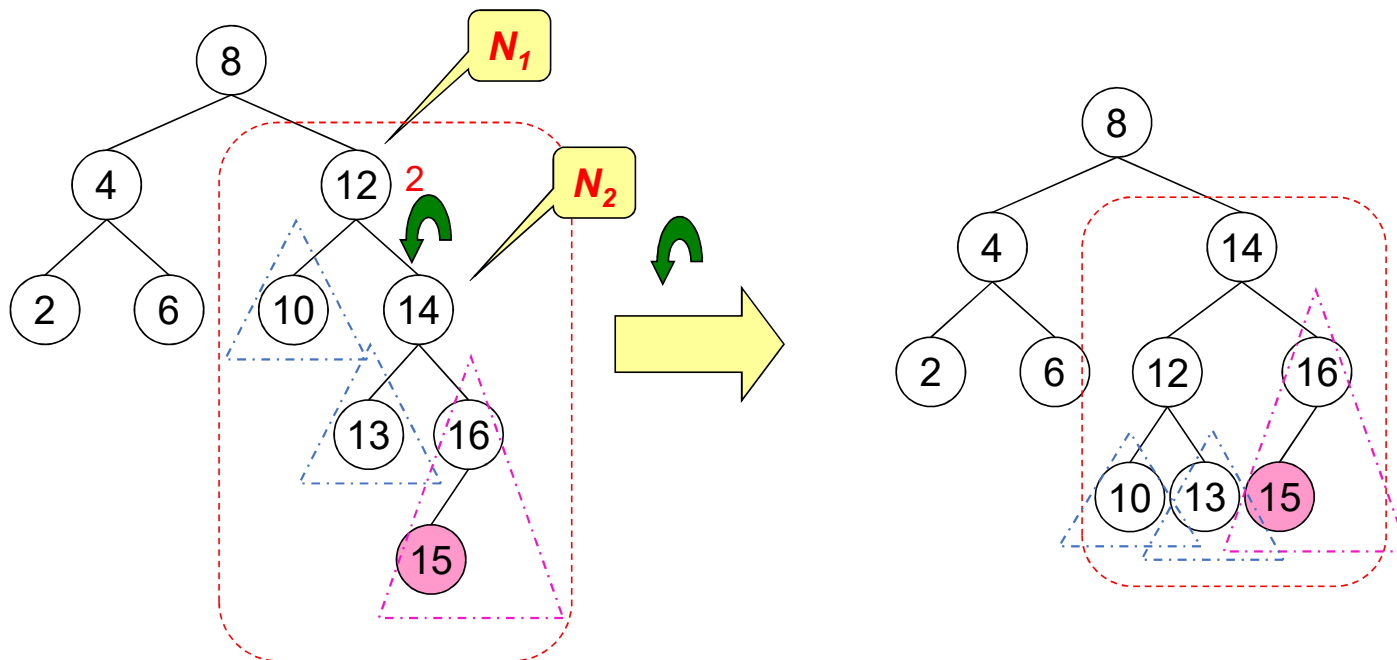



Left Rotation

- Key **15** is inserted to the **right** subtree of the **right** child (of node **12**). A **left** rotation is required. (Case 1)




Left Rotation

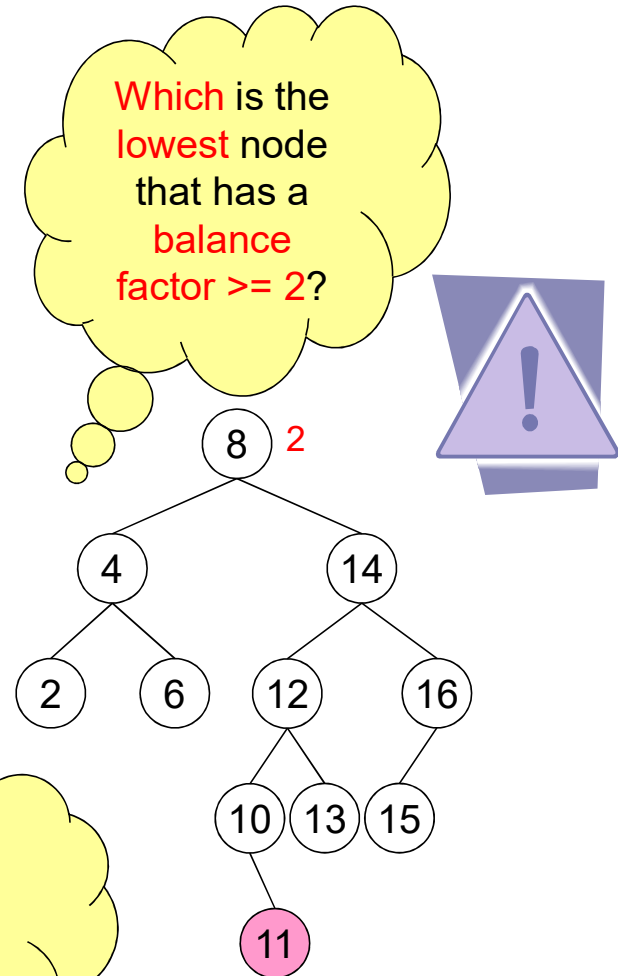



Example

- Finally, insert key 11.

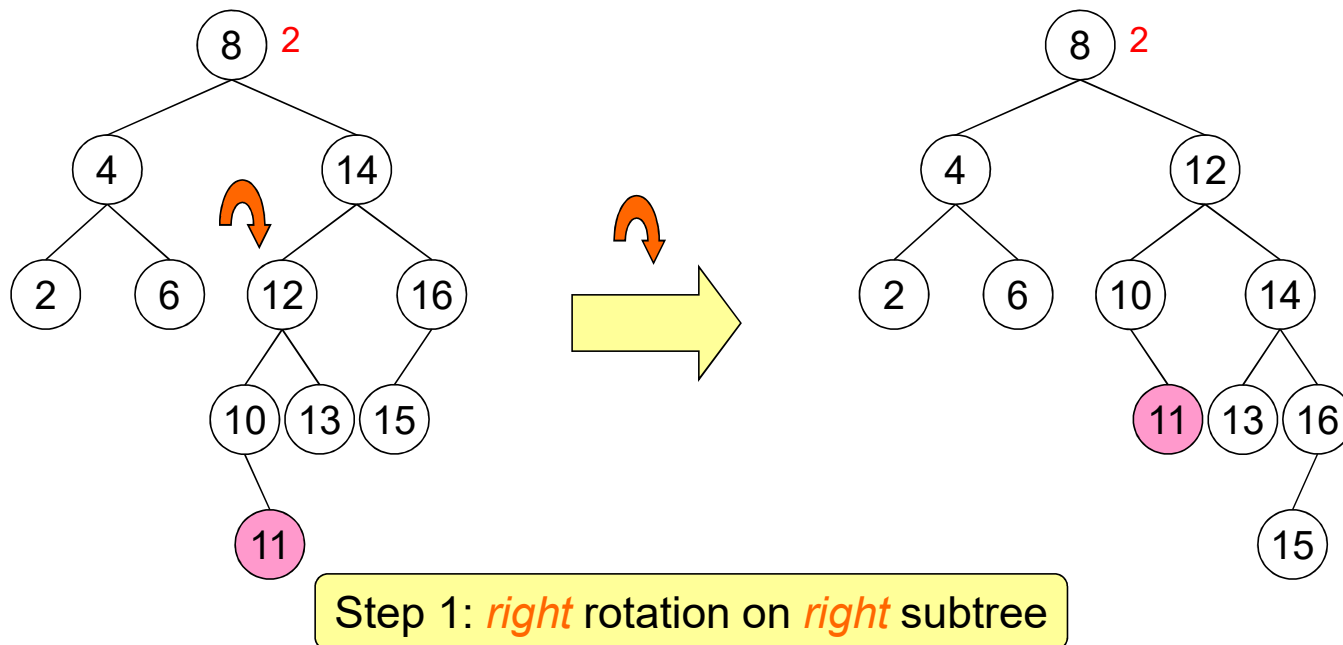


Insert 11



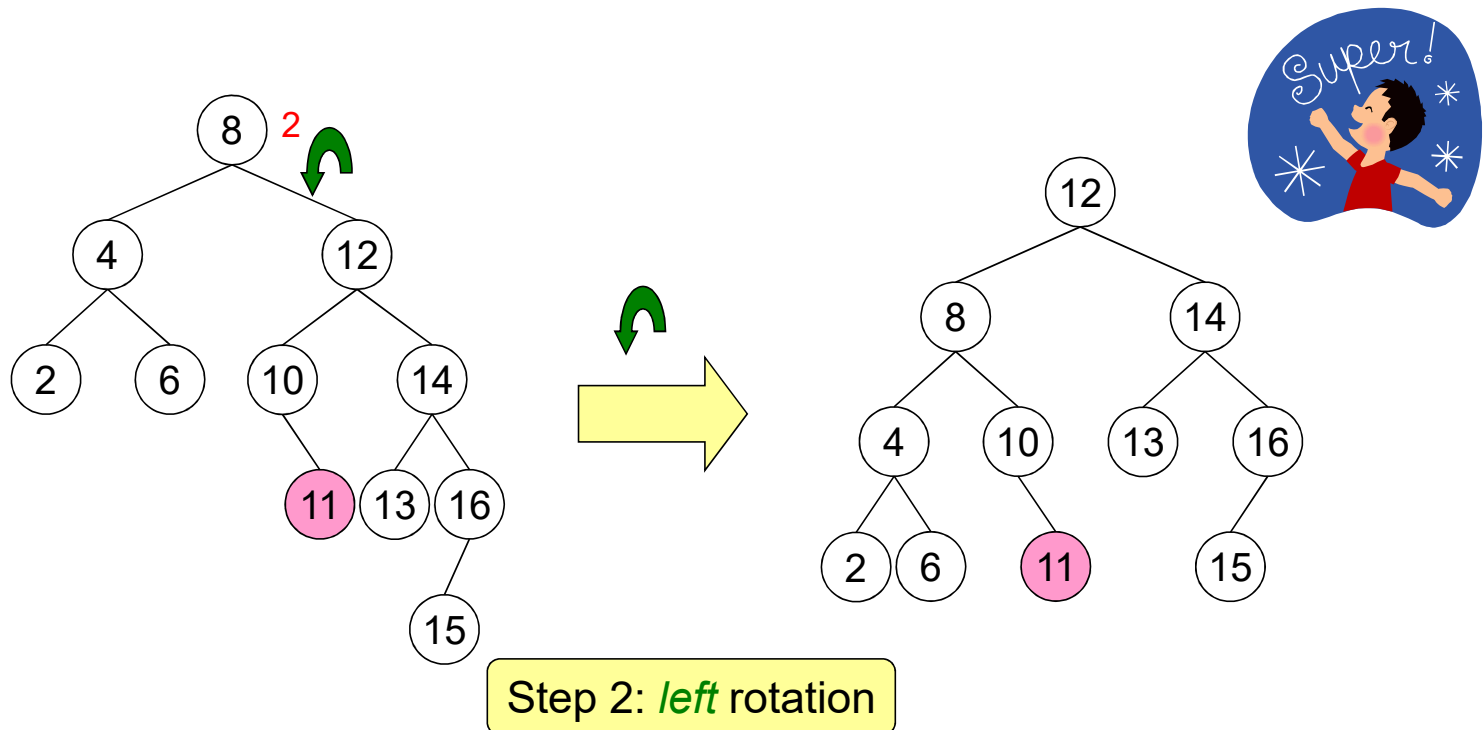
Right-Left Rotation



- Key **11** is inserted to the *left* subtree of the *right* child. A *right-left* rotation is required. (Case 4)



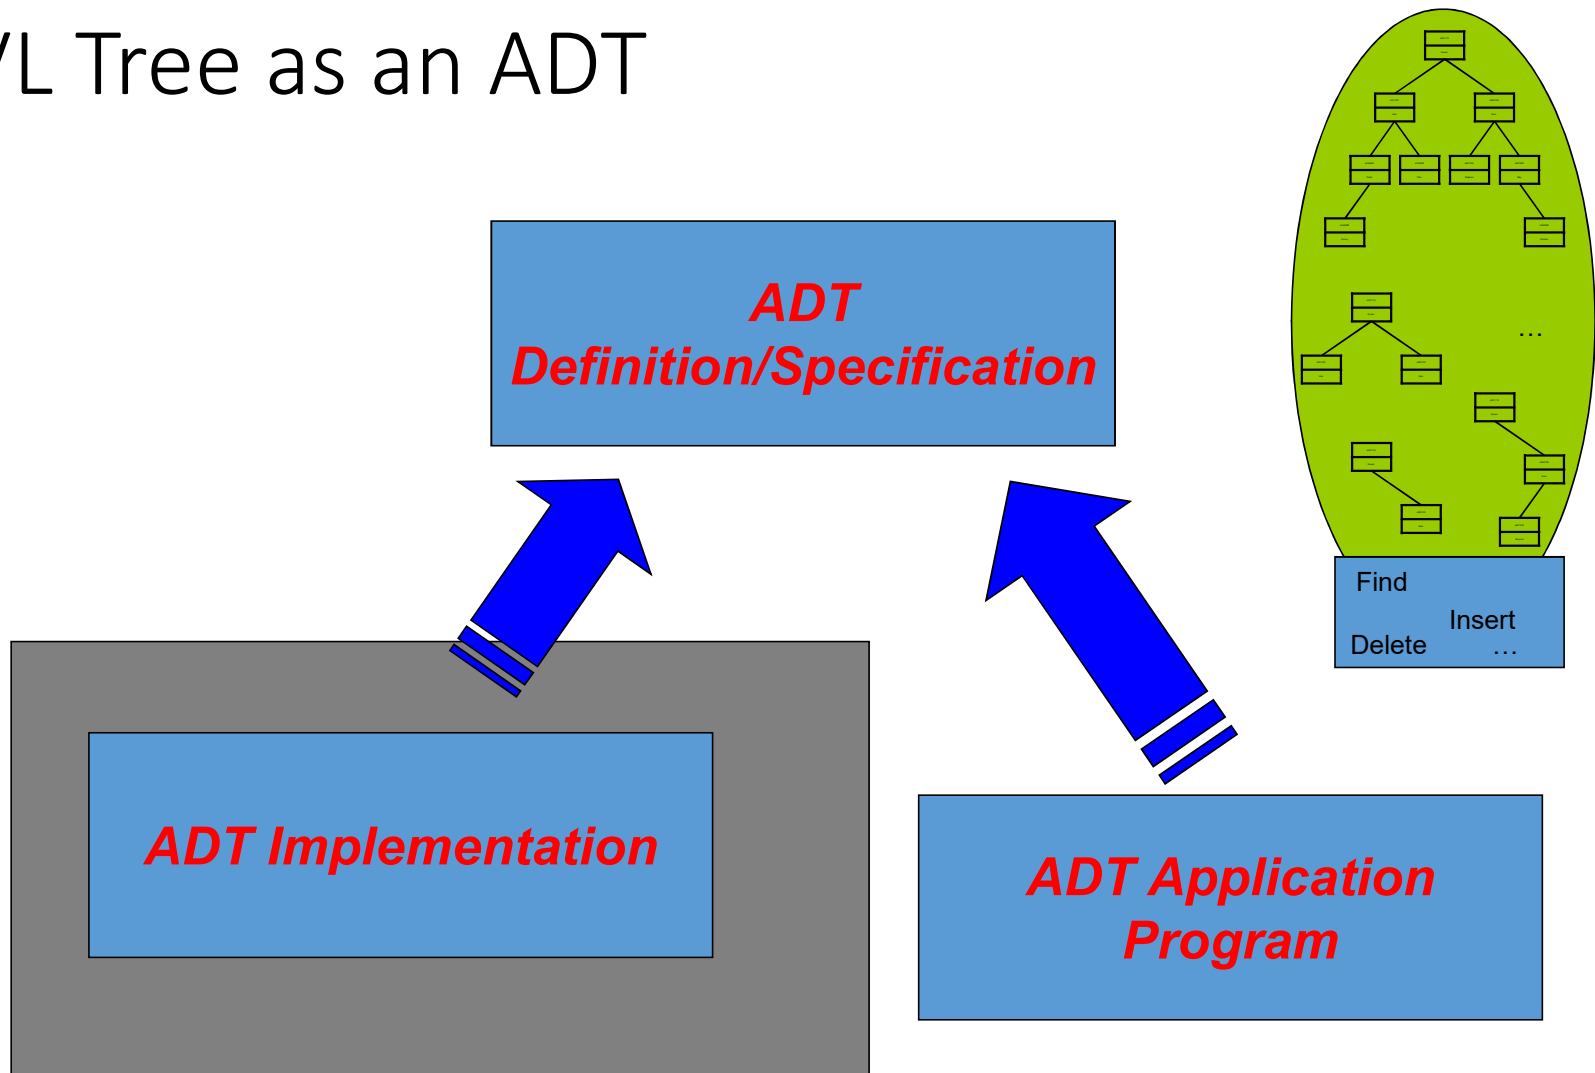
Right-Left Rotation

- Key **11** is inserted to the **left** subtree of the **right** child. A **right-left** rotation is required. (Case 4)





AVL Tree as an ADT



Reusing the BST ADT Interface

- AVL Trees can be seen as an extension to BSTs. We can *reuse* the interface of the BST ADT.

bst.h

```
#include "treenode.h"

typedef struct bstCDT *bstADT;

bstADT EmptyBST();
int BSTIsEmpty(bstADT t);

bstADT MakeBST(treeNodeADT root, bstADT left, bstADT right);

treeNodeADT Root(bstADT t);
bstADT LeftSubtree(bstADT t);
bstADT RightSubtree(bstADT t);

treeNodeADT FindNode(bstADT t, char *key);
bstADT InsertNode(bstADT t, treeNodeADT n);
bstADT DeleteNode(bstADT t, char *key);
```



Implementation of AVL Trees

- We modify the representation to explicitly store the *height* of a BST.

avltree.c

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "bst.h"

struct bstCDT {
    treeNodeADT root;
    bstADT left;
    bstADT right;
    int height;
};
```



avltree.c (continue)

```
int Height(bstADT t) {    // a private function
    if (BSTIsEmpty(t))
        return -1;
    else
        return t->height;
}

bstADT MakeBST(treeNodeADT root, bstADT left, bstADT right) {
    bstADT t;
    int lh, rh;
    t = (bstADT)malloc(sizeof(struct bstCDT));
    t->root = root;
    t->left = left;
    t->right = right;
    lh = Height(left);
    rh = Height(right);
    if (lh < rh)
        t->height = rh + 1;
    else
        t->height = lh + 1;
    return t;
}
```

Implementation of AVL Trees

- The functions **EmptyBST**, **BSTIsEmpty**, **Root**, **LeftSubtree**, **RightSubtree**, and **FindNode** are exactly the *same* as those in **bst.c**.

avltree.c (continue)

```
bstADT EmptyBST() { ... }

int BSTIsEmpty(bstADT t) { ... }

treeNodeADT Root(bstADT t) { ... }

bstADT LeftSubtree(bstADT t) { ... }

bstADT RightSubtree(bstADT t) { ... }

treeNodeADT FindNode(bstADT t, char *key) { ... }
```


Rotation Functions

- Next, we have to write (private) functions for the rotation operations.
 - Single Rotations
 - `bstADT RotateLeft(bstADT t)`
 - `bstADT RotateRight(bstADT t)`
 - Double Rotations
 - `bstADT RotateLeftRight(bstADT t)`
 - `bstADT RotateRightLeft(bstADT t)`

Single Rotation: RotateLeft

avltree.c (continue)

```
bstADT RotateLeft(bstADT t) {  
    treeNodeADT newRoot;  
    bstADT newLeft, newRight;  
    newRoot = Root(RightSubtree(t));  
    newLeft = MakeBST(Root(t), LeftSubtree(t),  
                      LeftSubtree(RightSubtree(t)));  
    newRight = RightSubtree(RightSubtree(t));  
    return MakeBST(newRoot, newLeft, newRight);  
}
```




Single Rotation: **RotateRight**

avltree.c (continue)

```
bstADT RotateRight(bstADT t) {
```

*(Entirely symmetrical to **RotateLeft**.)*
(Leave as an exercise.)

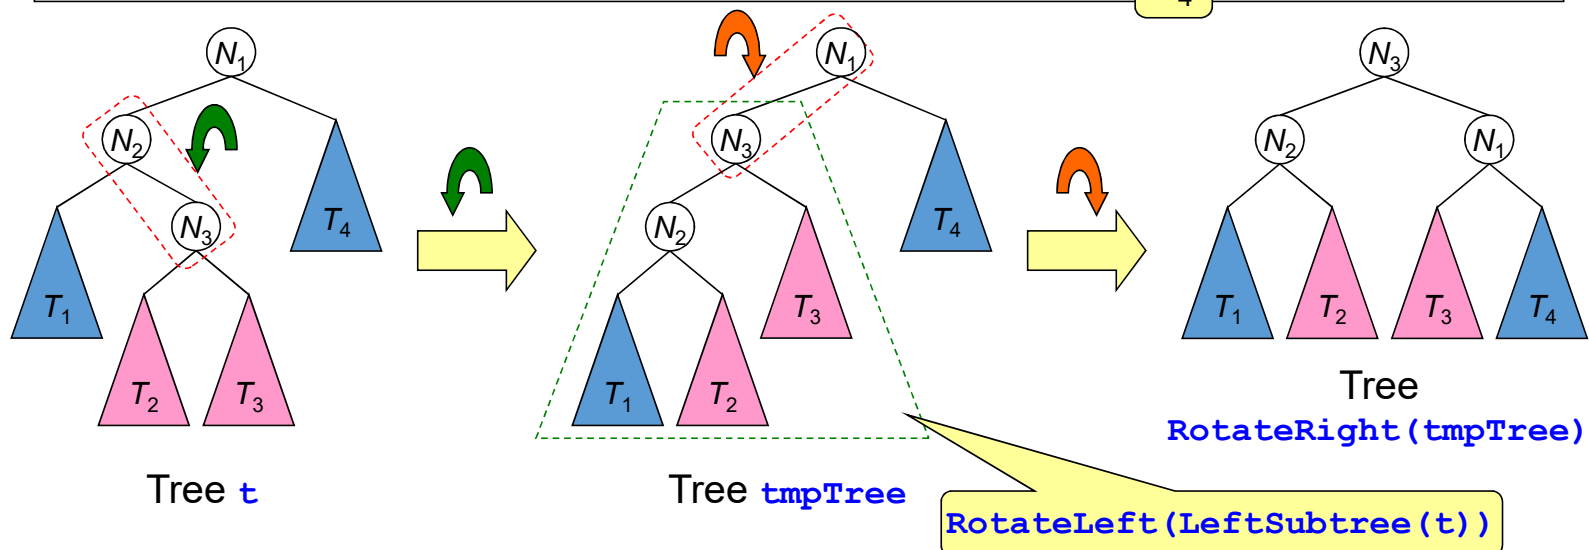
```
}
```



Double Rotation: RotateLeftRight

avltree.c (continue)

```
bstADT RotateLeftRight(bstADT t) {
    bstADT tmpTree;
    tmpTree = MakeBST(Root(t), RotateLeft(LeftSubtree(t)),
        N1 RightSubtree(t));
    return RotateRight(tmpTree);
}
```



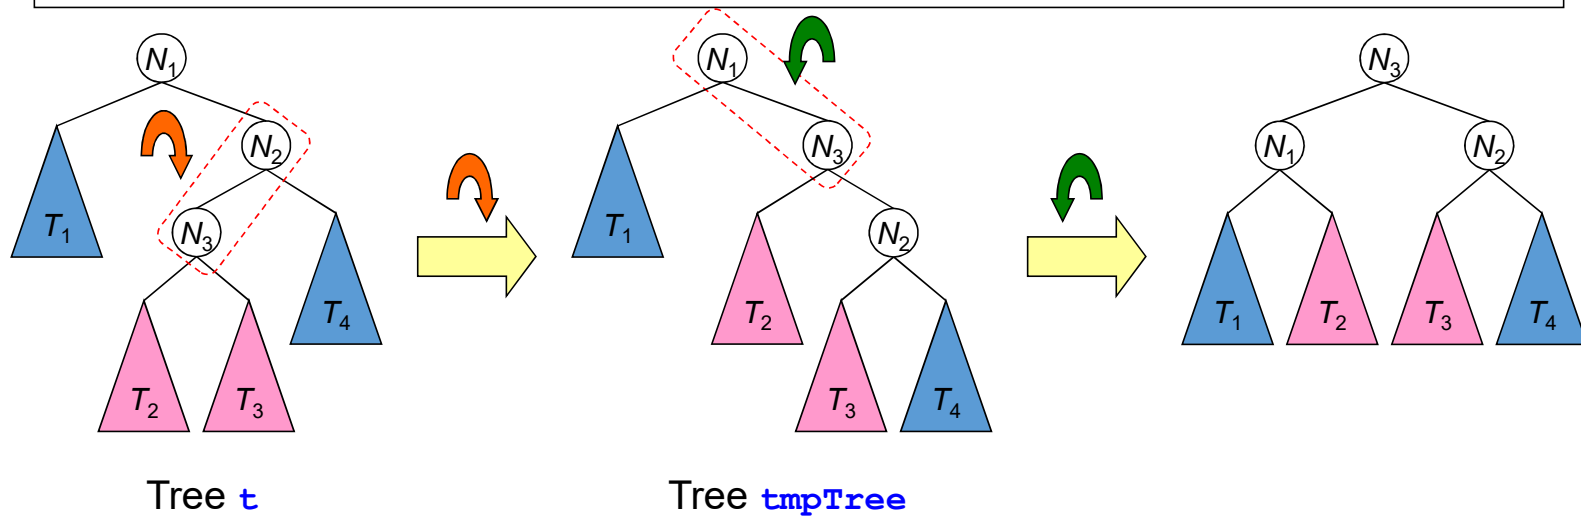
Double Rotation: **RotateRightLeft**

avltree.c (continue)

```
bstADT RotateRightLeft(bstADT t) {
```

*(Entirely symmetrical to **RotateLeftRight**.)*
(Leave as an exercise.)

```
}
```



AVL Node Insertion

- All the rotation functions are very easy and logical to understand.
- They all execute in $O(1)$ time.
- With the rotation functions, we are now ready to implement the **InsertNode** function, which is itself also very easy and logical.

AVL Node Insertion

avltree.c (continue)

```
bstADT InsertNode(bstADT t, treeNodeADT n) {
    int sign;
    bstADT tmpTree;
    if (BSTIsEmpty(t))
        return MakeBST(n, EmptyBST(), EmptyBST());
    sign = strcmp(GetNodeKey(n), GetNodeKey(Root(t)));
    if (sign == 0)
        return MakeBST(n, LeftSubtree(t), RightSubtree(t));
    else if (sign < 0) {    // insert to left subtree
        (See page 67)
    } else {    // insert to right subtree
        (See page 70)
    }
}
```

Same cases
as in [bst.c](#).

avltree.c (continue)


```
...
if (sign == 0)
    ...
else if (sign < 0) {    // insert to left subtree
    tmpTree = MakeBST(Root(t), InsertNode(LeftSubtree(t), n),
                      RightSubtree(t));

    if (Height(LeftSubtree(tmpTree))
        - Height(RightSubtree(tmpTree)) == 2) {
        if (strcmp(GetNodeKey(n),
                    GetNodeKey(Root(LeftSubtree(tmpTree)))) < 0)
            return RotateRight(tmpTree);    // case 2
        else
            return RotateLeftRight(tmpTree); // case 3
    } else
        return tmpTree;
} else {    // insert to right subtree
    ...
}
```

Check if the insertion causes imbalance.

tmpTree is balanced, simply return it.

If tmpTree is unbalanced, then determine whether it is case 2 or case 3.



```

if (Height(LeftSubtree(tmpTree))
    - Height(RightSubtree(tmpTree)) == 2) {
    if (strcmp(GetNodeKey(n),
        GetNodeKey(Root(LeftSubtree(tmpTree)))) < 0)
        return RotateRight(tmpTree);    // case 2
    else
        return RotateLeftRight(tmpTree); // case 3
} else
    return tmpTree;

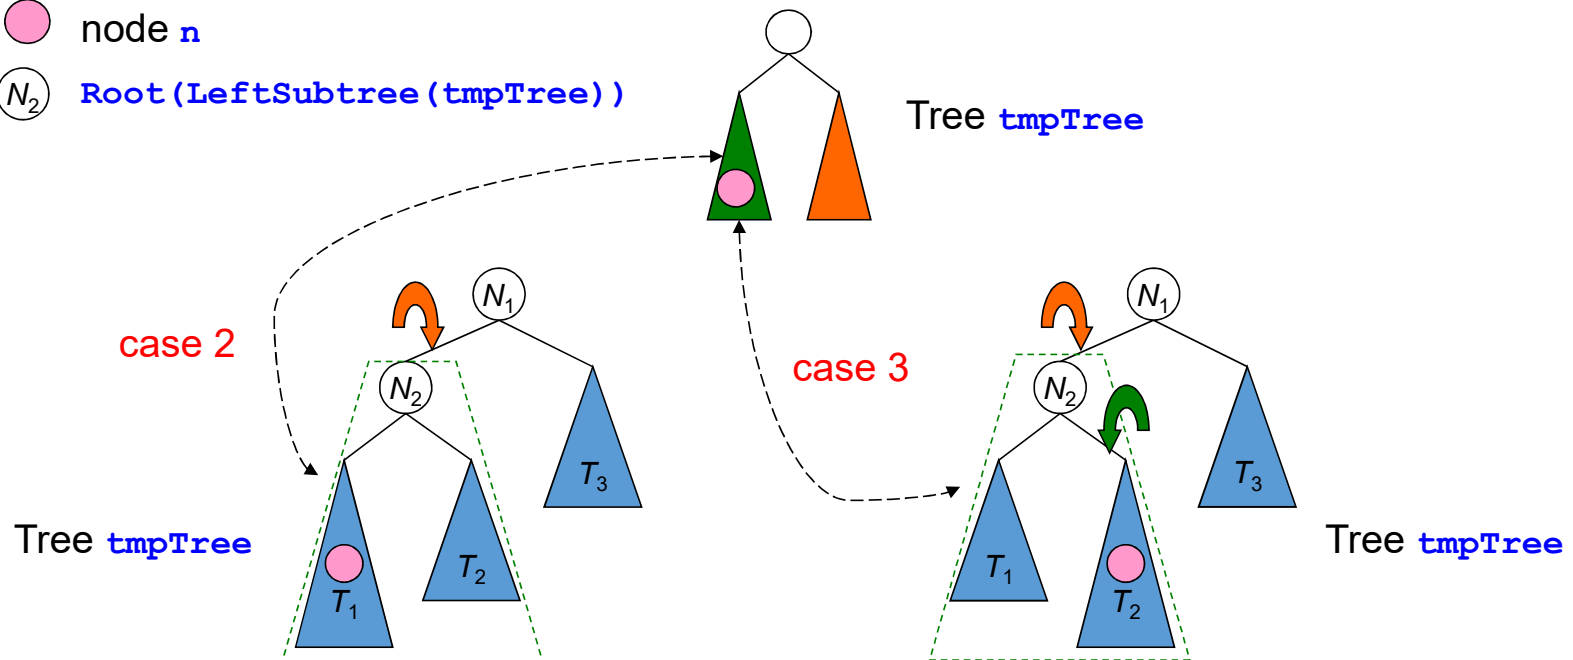
```



node n



N_2 $\text{Root}(\text{LeftSubtree}(\text{tmpTree}))$



AVL Node Insertion

avltree.c (continue)

```
bstADT InsertNode(bstADT t, treeNodeADT n) {
    int sign;
    bstADT tmpTree;
    if (BSTIsEmpty(t))
        return MakeBST(n, EmptyBST(), EmptyBST());
    sign = strcmp(GetNodeKey(n), GetNodeKey(Root(t)));
    if (sign == 0)
        return MakeBST(n, LeftSubtree(t), RightSubtree(t));
    else if (sign < 0) {    // insert to left subtree
        (Shown in page 67)
    } else {    // insert to right subtree
        (Shown in page 70)
    }
}
```

avltree.c (continue)



```
...
if (sign == 0)
    ...
else if (sign < 0) {    // insert to left subtree
    ...
} else {    // insert to right subtree
    tmpTree = MakeBST(Root(t), LeftSubtree(t),
                      InsertNode(RightSubtree(t), n));
    if (Height(RightSubtree(tmpTree))
        - Height(LeftSubtree(tmpTree)) == 2) {
        if (strcmp(GetNodeKey(n),
                    GetNodeKey(Root(RightSubtree(tmpTree)))) > 0)
            return RotateLeft(tmpTree);    // case 1
        else
            return RotateRightLeft(tmpTree);    // case 4
    } else
        return tmpTree;
}
```

Check if the insertion causes imbalance.

`tmpTree` is balanced, simply return it.

If `tmpTree` is unbalanced, then determine whether it is **case 1** or **case 4**.

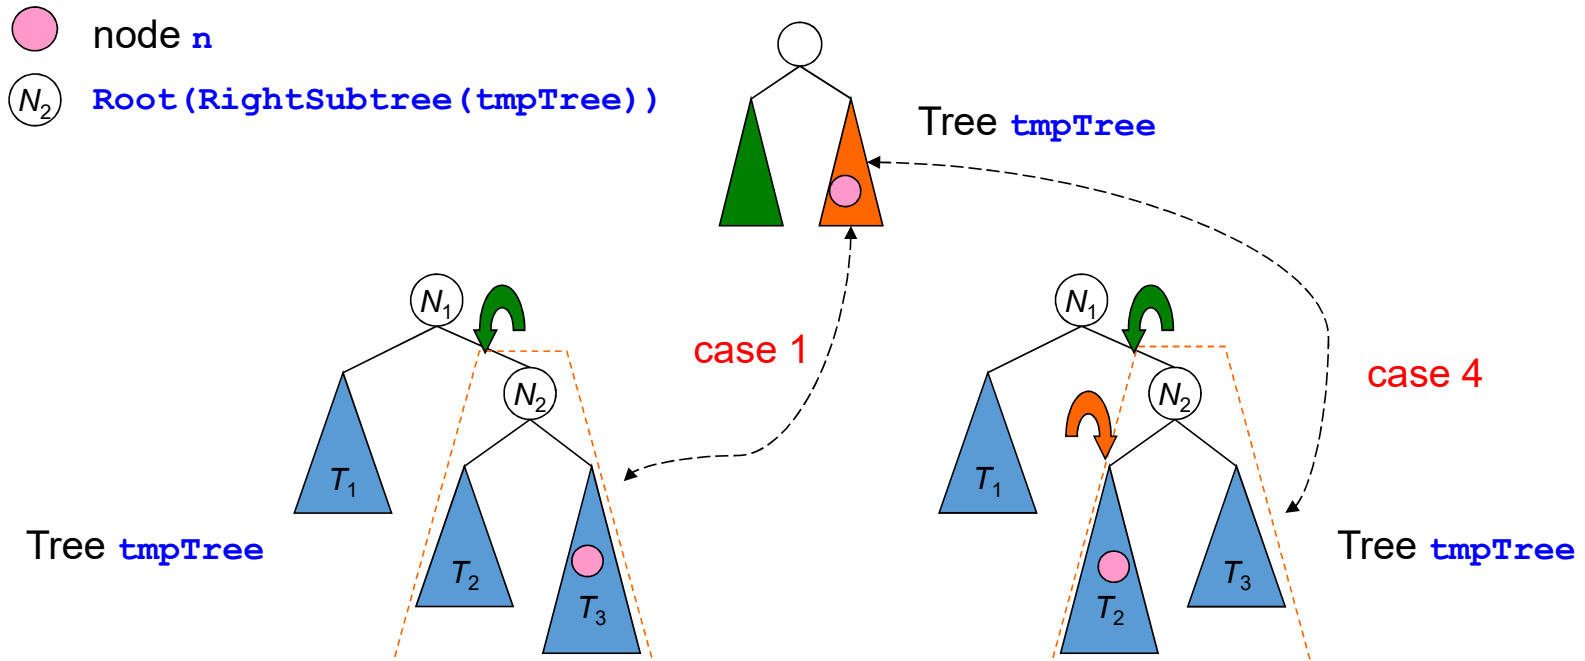
node `n`



```

if (Height(RightSubtree(tmpTree))
    - Height(LeftSubtree(tmpTree)) == 2) {
    if (strcmp(GetNodeKey(n),
                GetNodeKey(Root(RightSubtree(tmpTree)))) > 0)
        return RotateLeft(tmpTree);    // case 1
    else
        return RotateRightLeft(tmpTree); // case 4
} else
    return tmpTree;

```





Why lowest node?

- Find the lowest node that is imbalanced, i.e., the lowest node that has a balanced factor ≤ -2 or ≥ 2 , after an insertion
- Can we perform the rotation at any node that has a balanced factor ≤ -2 or ≥ 2 , after an insertion?

Consider this case

- Let's insert **11**, and perform the rotation at node **13** and **10**, respectively
- Do we get a balanced tree for both cases?




Correctness of **InsertNode**

- Does **InsertNode** find the lowest node that is imbalanced, i.e., the lowest node that has a balanced factor ≤ -2 or ≥ 2 , after an insertion?

Deletion in AVL Trees

- Deleting a node from an AVL tree can also cause imbalance, which has to be fixed also.



- The deletion algorithm will not be covered here.